

Master DevOps Interview Question Bank

1. Project / Resume / DevOps Experience

1. Tell me about yourself.
 2. Walk me through your current project architecture and your role in it.
 3. Can you walk me through your previous project architecture and your roles and responsibilities?
 4. Which DevOps tools have you worked with in the last 2 years?
 5. What AWS services have you used in production?
 6. Explain your day-to-day activities.
 7. Do you have hands-on experience in Linux? If yes, which platform?
 8. Can you walk me through your CI process?
 9. Can you walk me through your deployment process?
 10. Describe your experience with CI/CD pipelines.
 11. What AWS services have you used in your project?
 12. Walk me through your current project architecture.
 13. What is your role and responsibility in the project?
 14. How do you handle production support and troubleshooting?
-

2. Linux

15. What are the most common Linux commands you use?
16. How do you check running processes in Linux?
17. How do you list running processes?
18. How do you check one particular running process?
19. How do you find a particular process ID?
20. How do you kill a process?
21. Can you find a process ID and kill it using a single command?
22. How do you check disk usage?
23. How do you check free memory?

24. How do you check CPU utilization?
25. How do you check swap usage?
26. What does swapon --show do?
27. How do you check filesystem usage?
28. How do you find files larger than 100 MB?
29. How do you find the top largest files in a directory?
30. How do you archive and compress a directory in Linux?
31. What does chmod 755 mean?
32. If you execute chmod 755 <file>, what exactly happens?
33. What is shown by ls -l?
34. How do you list all SSH users in Linux?
35. How do you check which users are currently logged in?
36. How do you check SSH login history?
37. If an Apache server is running and you need to check its logs, where do you look?
38. In which directory are Apache logs generally stored?
39. What kind of logs can you see under /var/log?
40. What is the purpose of the grep command?
41. If you want to grep the keyword linux from linux.txt, which flag would you use to list all matching lines?
42. How do you search for a particular string across multiple files?
43. How do you search recursively using grep?
44. How do you check which process is listening on a particular port?
45. How do you check whether a service is listening on port 80?
46. How do you check network connectivity between two servers?
47. What commands would you use to troubleshoot network connectivity?
48. How do you check whether a port is reachable from another server?
49. How do you identify which process is using a particular port?
50. What is lsof?

51. How do you use lsof to find which process is using a port?
 52. What is the difference between systemctl and journalctl?
 53. What is systemd?
 54. What is journald?
 55. Where are application logs stored?
 56. How are system/service logs different from application logs?
 57. How do you check the status of a systemd service?
 58. How do you start, stop and restart a service?
 59. What does systemctl cat nginx do?
 60. How do you troubleshoot a service that is running but not accessible?
 61. How do you troubleshoot Apache/Nginx if it is listening on the expected port but users cannot access it?
 62. How do you check firewall rules in Linux?
 63. How do you identify whether a firewall is blocking traffic?
 64. What happens if a file is being used by another application? Can you delete it?
 65. How do you find which process has a file open?
 66. What would you do if an application is holding a file open?
-

3. Linux Scheduling

67. What is the difference between cron and at?
 68. How do you schedule an activity in Linux?
 69. How do you run a script every 5 minutes?
 70. How do you configure a cron job?
 71. Where is the crontab configuration stored?
 72. How do you check existing cron jobs?
 73. How do you troubleshoot a cron job that is not executing?
-

4. Bash Scripting

- 74. Can you write a Bash script to find the top 5 largest files in a directory and display their sizes?
 - 75. How would you write a Bash script to perform a node-health check?
 - 76. What checks would you include in a node-health Bash script?
 - 77. How would you check CPU utilization from Bash?
 - 78. How would you check memory utilization from Bash?
 - 79. How would you check disk utilization from Bash?
 - 80. How would you check whether a particular service is running from Bash?
 - 81. How would you check whether a particular port is listening using Bash?
 - 82. How would you check whether a server is reachable using Bash?
 - 83. How would you write a script that alerts when disk usage exceeds 80%?
 - 84. How would you write a script to find failed services?
 - 85. How would you write a script to monitor a process?
 - 86. How would you write a script to find the most common IP address in a log file?
 - 87. How would you write a Bash script to parse application logs?
 - 88. What error handling mechanisms do you use in Bash scripts?
-

5. Python Scripting

- 89. Write a Python script to read a log file and count how many times the word ERROR appears.
 - 90. Could you write a Python script that backs up all files older than 30 days from a folder?
 - 91. How would you find files older than 30 days using Python?
 - 92. How would you handle exceptions in a Python automation script?
 - 93. How would you make a Python automation script accept command-line arguments?
-

6. Git

- 94. Explain git pull and git fetch.

- 95. What is the difference between git pull and git fetch?
 - 96. How do you resolve a Git merge conflict?
 - 97. What is the difference between git merge and git rebase?
 - 98. How do you resolve merge conflicts during a rebase?
 - 99. What is git cherry-pick?
 - 100. What is git config?
 - 101. What is git rebase?
 - 102. How do you undo the last Git commit?
 - 103. What is the difference between git reset and git revert?
 - 104. How do you check Git branches?
 - 105. How do you create and delete a Git branch?
 - 106. How do you push a local branch to a remote repository?
 - 107. How do you handle merge conflicts in a large repository?
-

7. GitHub / GitHub Actions

- 108. How do you generate a token in GitHub?
 - 109. Explain the steps to generate a GitHub token.
 - 110. How does a GitHub Actions workflow get triggered?
 - 111. Explain the structure of a GitHub Actions YAML file.
 - 112. What are jobs and steps in GitHub Actions?
 - 113. How do you pass secrets to GitHub Actions?
 - 114. How do you configure environment-specific deployments in GitHub Actions?
 - 115. How do you implement CI/CD using GitHub Actions?
-

8. GitLab CI/CD

- 116. What is the difference between GitHub Actions and GitLab CI/CD?
- 117. What are stages and jobs in GitLab pipelines?

- 118. Explain the structure of a .gitlab-ci.yml file.
 - 119. What are stages, jobs and pipelines in GitLab?
 - 120. What is the difference between artifacts and cache in GitLab?
 - 121. How do GitLab runners work?
 - 122. What is a protected branch?
 - 123. What is a protected variable?
 - 124. How do you trigger pipelines manually in GitLab?
 - 125. How do you configure environment-specific deployments in GitLab CI/CD?
-

9. Jenkins Fundamentals

- 126. Explain Jenkins architecture.
- 127. What is a Jenkins pipeline?
- 128. What are Jenkins pipeline stages?
- 129. What is the difference between Declarative and Scripted Jenkins pipelines?
- 130. Can you write a complete Jenkins pipeline?
- 131. How do you implement CI/CD for a containerized application?
- 132. How do you secure Jenkins credentials?
- 133. How do you install Jenkins plugins?
- 134. What are the steps to install a Jenkins plugin?
- 135. If Jenkins is installed but not accessible, what could be the reason?
- 136. How do you troubleshoot Jenkins when the service is running but the UI is inaccessible?
- 137. How do you integrate SonarQube with Jenkins?
- 138. How do you integrate Maven with Jenkins?
- 139. How do you configure Jenkins to build an application?
- 140. What is a Jenkins agent?
- 141. What is the difference between Jenkins controller and agent?

- 142. How do you add a Linux server as a Jenkins worker node?
 - 143. If someone gives you a Linux server and asks you to add it as a Jenkins worker node, what steps would you follow?
 - 144. How do you configure Jenkins agents securely?
 - 145. How do you handle a Jenkins agent becoming unresponsive during deployment?
 - 146. How would you make a Jenkins pipeline self-healing if a shared agent becomes unresponsive?
 - 147. How do you implement error handling in Jenkins?
 - 148. How do you implement retry logic in Jenkins?
 - 149. How do you configure Jenkins pipeline notifications?
 - 150. How do you prevent multiple conflicting Jenkins builds from running simultaneously?
-

10. Jenkins / Maven / Java

- 151. What is the Maven lifecycle?
 - 152. What is the mvn clean install command?
 - 153. What is a WAR file?
 - 154. Where do you store the generated artifact?
 - 155. Where can you find the generated artifact file?
 - 156. What happens after running a Node.js application?
 - 157. How do you configure multiple Java versions on a Linux Jenkins worker?
 - 158. If a Linux worker has multiple Java versions, how would you write a Jenkins pipeline to use them?
 - 159. What do you write inside a Jenkins pipeline to make use of multiple Java versions?
 - 160. How do you configure different Java versions for different Jenkins jobs?
-

11. CI/CD — Advanced Scenarios

161. How do you design a controlled deployment flow where every release must pass environment-specific validation before progressing?
 162. What is your approach to implementing artifact promotion across dev → QA → staging → prod without rebuilding?
 163. How would you ensure pipeline self-healing when a shared agent becomes unresponsive mid-deployment?
 164. How do you handle parallel infrastructure changes in a monorepo environment to avoid conflicting pipeline triggers?
 165. How do you design a CI/CD pipeline that supports multiple release calendars for different business units?
 166. What is your strategy for enforcing branch protection and release gates in large repositories?
 167. How do you handle failed deployments when rollback scripts are not available?
 168. How do you implement artifact versioning and promotion across environments?
 169. How do you design a CI/CD pipeline where failed deployments must auto-freeze further releases until the root cause is identified?
 170. How would you implement deployment verification tests for systems that cannot tolerate downtime?
 171. Explain how you would manage version pinning across multiple dependent microservices.
 172. How do you ensure CI pipelines remain deterministic when builds depend on external services?
 173. How do you verify that a deployment actually changed runtime behavior?
 174. How do you debug a pipeline that succeeds but produces a faulty release?
 175. How do you implement approval gates in a production deployment pipeline?
 176. How do you implement automated rollback in CI/CD?
 177. How do you promote the same artifact from development to production without rebuilding it?
-

12. Docker Fundamentals

178. What is Docker?
179. What is the difference between a Docker image and a container?
180. How do you create a custom Docker image?
181. What is a Dockerfile?
182. Explain a Dockerfile.
183. What are the prerequisites for the docker build command?
184. What happens when you run docker build?
185. What is the difference between COPY and ADD?
186. When would you prefer COPY over ADD?
187. What are CMD and ENTRYPOINT?
188. What is the difference between CMD and ENTRYPOINT?
189. What are Docker volumes and why are they used?
190. What is the difference between a Dockerfile and Docker Compose?
191. How do you copy a file from a container to the host?
192. If we write COPY . . , what happens in the background?
193. What is the source and destination when using COPY . . ?
194. What is a Docker build context?
195. What is a Docker image layer?
196. How does Docker caching work?
197. How do you reduce Docker image size?
198. What is a multi-stage Docker build?
199. Why are multi-stage Docker builds used?
200. Write a Dockerfile for a Node.js application using a multi-stage build.
201. Can you write a practical Node.js Dockerfile that you would actually use in a real project?
202. How do you troubleshoot a container that has exited?
203. A container has crashed or exited suddenly. How would you determine what went wrong?

13. Docker Security / Application Secrets

- 204. How would you pass MySQL credentials securely to a PHP application running in a Docker container?
- 205. How would you handle secrets in a Docker container for a PHP application connecting to MySQL?
- 206. How do you avoid hardcoding secrets in Dockerfiles?
- 207. How do you securely inject runtime configuration into containers?
- 208. How do you scan Docker images for vulnerabilities?
- 209. How do you prevent vulnerable images from reaching production?

14. Kubernetes Fundamentals

- 210. What is Kubernetes?
- 211. What is a Kubernetes cluster?
- 212. What is a Pod?
- 213. What is a Node?
- 214. What is a Deployment?
- 215. What is a ReplicaSet?
- 216. What is a StatefulSet?
- 217. What is a DaemonSet?
- 218. What is a Job?
- 219. What is a CronJob?
- 220. What is a Service?
- 221. What is a Namespace?
- 222. What is ConfigMap?
- 223. What is a Secret?
- 224. What is a Volume?
- 225. What is a PersistentVolume?
- 226. What is a PersistentVolumeClaim?

- 227. What is a StorageClass?
 - 228. What is Kubelet?
 - 229. What is Kube-proxy?
 - 230. What is the API Server?
 - 231. What is the Controller Manager?
 - 232. What is the Scheduler?
 - 233. What is the difference between Deployment, StatefulSet and DaemonSet?
 - 234. How does Kubernetes perform rolling updates?
 - 235. How does Kubernetes handle pod restarts and failures?
 - 236. What happens when a Kubernetes node goes into NotReady state?
 - 237. What is the difference between ClusterIP, NodePort and LoadBalancer services?
 - 238. How do liveness and readiness probes work?
 - 239. What is the purpose of ConfigMaps and Secrets?
 - 240. How does the Horizontal Pod Autoscaler work internally?
 - 241. What is the difference between Ingress and Service?
 - 242. What is the use of namespaces in Kubernetes?
 - 243. How do you troubleshoot a pod stuck in CrashLoopBackOff?
-

15. Kubernetes Networking

- 244. How do Kubernetes Services find Pods?
- 245. How are Kubernetes Services linked to Pods?
- 246. How does a Service select the correct Pods?
- 247. How does one namespace Pod communicate with a Pod in another namespace?
- 248. How does Kubernetes DNS work?
- 249. What DNS name can a Pod use to access a Service in another namespace?

- 250. How do you expose a Kubernetes application to external traffic?
 - 251. What is the difference between ClusterIP, NodePort and LoadBalancer?
 - 252. What is the difference between Ingress and Service?
 - 253. How would you troubleshoot a Kubernetes Pod that is running but traffic is failing?
 - 254. What do you check first when a Kubernetes Pod is running but traffic fails?
 - 255. How do you detect and fix uneven Pod traffic distribution behind a Service?
 - 256. How do you implement NetworkPolicies to restrict Pod-to-Pod communication?
 - 257. How do you enforce rules in Kubernetes to control which Pods can talk to each other?
 - 258. How would you validate the impact of a new NetworkPolicy before applying it cluster-wide?
-

16. Kubernetes Deployments / Application Operations

- 259. How do you configure a rolling update in Kubernetes?
- 260. How do you perform a rolling update?
- 261. How do you implement a blue-green deployment in Kubernetes?
- 262. How would you roll out a blue-green deployment in a Kubernetes cluster?
- 263. What would a blue-green deployment look like in practice?
- 264. How do you configure deployment strategies in Kubernetes?
- 265. How do you manage ConfigMap changes without restarting applications?
- 266. How do you safely rotate TLS certificates used by Ingress controllers without service disruption?
- 267. How do you troubleshoot Pods stuck in Terminating state in production?
- 268. How do you isolate workloads of different teams inside a shared Kubernetes cluster?
- 269. How do you isolate workloads using namespaces?

- 270. How do you enforce namespace-level resource isolation in a multi-team Kubernetes cluster?
 - 271. How do you protect critical workloads from noisy-neighbor issues in shared clusters?
 - 272. What are resource requests and limits and why are they important?
 - 273. How do you safely roll out kernel or node-image upgrades in production Kubernetes clusters?
-

17. Kubernetes Scheduling / Nodes / Runtime Reliability

- 274. A StatefulSet Pod is failing only on a specific node. Walk through how you would determine whether it is a node issue, storage issue or scheduling conflict.
- 275. What steps do you take if kubelet health degrades on only a subset of nodes?
- 276. How do you debug sporadic Pod restarts when logs show no application errors?
- 277. How do you troubleshoot a Pod that is running on one node but fails on another?
- 278. How do you investigate scheduling conflicts?
- 279. How do you troubleshoot Pods stuck in Pending state?
- 280. How do you troubleshoot Pods stuck in ImagePullBackOff?
- 281. How do you troubleshoot Pods stuck in ErrImagePull?
- 282. How do you troubleshoot Pod eviction?
- 283. What causes Kubernetes Pod eviction?
- 284. How do you troubleshoot node resource pressure?
- 285. How do you troubleshoot memory pressure on Kubernetes nodes?
- 286. How do you troubleshoot disk pressure on Kubernetes nodes?
- 287. How do you safely drain a Kubernetes node?
- 288. How do you determine whether a problem is application-level or node-level?
- 289. How do you troubleshoot intermittent Pod failures that cannot be reproduced?

18. Kubernetes Storage

- 290. Explain the relationship between Pod, Volume, PVC and PV.
- 291. What is a PersistentVolume?
- 292. What is a PersistentVolumeClaim?
- 293. What is a StorageClass?
- 294. How does dynamic volume provisioning work?
- 295. What is the role of the EBS CSI driver?
- 296. What is the role of the EFS CSI driver?
- 297. What is the difference between EBS and EFS?
- 298. What is gp3?
- 299. What does fsType: ext4 mean?
- 300. What are Kubernetes access modes?
- 301. How would you troubleshoot a StatefulSet storage failure?
- 302. How would you determine whether a StatefulSet problem is related to storage or scheduling?

19. Kubernetes Autoscaling / Reliability

- 303. How does Kubernetes Horizontal Pod Autoscaler work?
- 304. How do you confirm whether autoscaling is working as expected?
- 305. What metrics do you use for HPA?
- 306. How do you troubleshoot an HPA that is not scaling?
- 307. How does Cluster Autoscaler differ from HPA?
- 308. How would you protect critical workloads from resource starvation?
- 309. How would you design Kubernetes workloads for high availability?
- 310. How do you ensure workloads are distributed across Availability Zones?

20. Kubernetes Production Troubleshooting

311. How do you debug a production issue when logs show no errors?
312. What steps do you follow when application latency suddenly increases?
313. How do you identify the root cause when metrics look normal but users complain?
314. What do you check first when a Kubernetes Pod is running but traffic fails?
315. How do you troubleshoot intermittent failures that cannot be reproduced?
316. How do you verify whether a deployment actually changed runtime behavior?
317. What do you analyze when database connections are exhausted but CPU is low?
318. How do you approach an incident involving multiple dependent services?
319. What signals do you trust most during a live production outage?
320. How do you confirm whether autoscaling is working as expected?
321. How do you debug a pipeline that succeeds but produces a faulty release?
322. What would you investigate if a rollback does not resolve the issue?
323. How do you detect configuration drift in production systems?
324. What checks do you perform before declaring an incident resolved?
325. How do you prioritize actions during a high-severity outage?
326. How do you distinguish application bugs from infrastructure issues during incidents?
327. What metrics do you prioritize when troubleshooting slow API responses?
328. How do you design alerting rules that adapt to business-hour versus off-hour traffic patterns?
329. How do you correlate logs, metrics and traces for faster root-cause analysis?
330. How do you distinguish application latency from infrastructure latency during incidents?
331. What signals would you monitor to detect silent failures?
332. How do you troubleshoot an application returning HTTP 503 errors?

- 333. How do you determine whether a 503 is caused by the application, Service, Ingress or Load Balancer?
 - 334. How do you troubleshoot a production Kubernetes cluster where Pods are stuck in ImagePullBackOff, some Pods are being evicted, and users are reporting 503 errors?
 - 335. How would you prevent the above Kubernetes issues from recurring?
-

21. ArgoCD / GitOps

- 336. How do you integrate ArgoCD for deployment?
- 337. Once the CI build is completed and the image is pushed to ECR, what configurations do you configure on top of ArgoCD to perform deployment?
- 338. How does ArgoCD connect to a Kubernetes cluster?
- 339. How does ArgoCD perform the deployment?
- 340. We have ArgoCD and ECR. Whenever you push an image to ECR, how does ArgoCD know that there is an image change?
- 341. How does ArgoCD know which Kubernetes cluster it needs to deploy to?
- 342. How does ArgoCD get to know that there is an image change and that there is a new image in ECR?
- 343. How do you integrate ECR and ArgoCD so that image changes trigger deployment?
- 344. When you open the ArgoCD dashboard, what do you have?
- 345. What do we call a Kubernetes cluster in ArgoCD?
- 346. How does ArgoCD maintain the desired state?
- 347. What is GitOps?
- 348. What is the difference between CI and GitOps-based CD?
- 349. What is an ArgoCD Application?
- 350. What is the difference between Synced and OutOfSync in ArgoCD?
- 351. How do you troubleshoot an ArgoCD application showing OutOfSync?
- 352. How do you configure ArgoCD for multiple environments?
- 353. How do you configure ArgoCD for multiple Kubernetes clusters?

- 354. How does ArgoCD detect changes in Git?
 - 355. How do you implement automated synchronization in ArgoCD?
 - 356. How do you implement manual approval before production synchronization?
-

22. Ansible

- 357. How do you configure rolling updates using an Ansible playbook?
- 358. I have 100 worker nodes. When I run the playbook, it should run on the first 20 servers, then the next 20, then the next 20. How do you configure this?
- 359. How do you implement batching in Ansible?
- 360. I have 100 servers and want to fetch the IP address and OS family of those servers using an Ansible playbook. How would you do it?
- 361. Do you know how to execute a specific task in Ansible?
- 362. How do you execute a specific task using Ansible?
- 363. When you run an Ansible playbook, what do you get at the end?
- 364. What is the Ansible playbook recap?
- 365. What is the folder structure of Ansible roles?
- 366. Why do we use roles in Ansible?
- 367. Can you explain the Ansible role folder structure?
- 368. If I give you one Linux server and ask you to add it as a Jenkins worker node, how can you do that using Ansible?
- 369. How do you manage multiple Java versions using Ansible?
- 370. How do you use Ansible inventory?
- 371. What is the difference between static and dynamic inventory?
- 372. What are Ansible facts?
- 373. How do you gather facts from remote servers?
- 374. How do you execute an Ansible task on a subset of servers?
- 375. How do you use serial in Ansible?
- 376. How do you make Ansible playbooks idempotent?

- 377. How do you handle failures in Ansible?
 - 378. How do you use handlers in Ansible?
 - 379. How do you pass variables to an Ansible playbook?
 - 380. What is the difference between group_vars and host_vars?
 - 381. How do you use Ansible Vault?
-

23. Terraform Fundamentals

- 382. What is Infrastructure as Code?
- 383. What is Terraform?
- 384. What is the Terraform state file?
- 385. What is the purpose of the Terraform state file?
- 386. What is the difference between terraform init, terraform plan and terraform apply?
- 387. What happens during terraform init?
- 388. What happens during terraform plan?
- 389. What happens during terraform apply?
- 390. How do you manage multiple environments in Terraform?
- 391. What are Terraform workspaces?
- 392. What are Terraform modules?
- 393. Why do we use Terraform modules?
- 394. How do you structure Terraform modules?
- 395. How do you manage Terraform variables?
- 396. What are Terraform outputs?
- 397. What is a Terraform backend?
- 398. Why do you use a remote backend?
- 399. How do you handle Terraform state locking?
- 400. How do you handle Terraform state file corruption?
- 401. What happens if the Terraform state file gets corrupted or becomes out of sync?

- 402. How do you recover from Terraform state corruption?
 - 403. How do you import an existing AWS VPC into Terraform?
 - 404. What are the steps to import an existing AWS resource into Terraform?
 - 405. How do you make sure importing an existing VPC does not break anything?
 - 406. How do you detect Terraform configuration drift?
 - 407. How do you handle infrastructure drift?
 - 408. How do you migrate manually created infrastructure into Terraform?
 - 409. How do you manage dependencies between Terraform resources?
 - 410. What is the Terraform resource dependency graph?
 - 411. What is the difference between implicit and explicit dependencies?
-

24. Terraform Advanced

- 412. How do you structure IaC for multi-account, multi-region deployments with shared resource modules?
- 413. What is your strategy for detecting stale or abandoned cloud resources created through automation?
- 414. How do you migrate IaC from an in-house script-based model to Terraform without breaking dependencies?
- 415. How do you manage secrets in Terraform without hardcoding them?
- 416. How do you securely pass sensitive values to Terraform?
- 417. How do you implement cross-account resource provisioning using Terraform?
- 418. How would you create resources across two different AWS accounts using Terraform?
- 419. If an S3 bucket was created through Terraform but someone manually added a policy, how would you handle the drift?
- 420. If someone manually changed an S3 bucket policy that was originally created with Terraform, how would you handle it?
- 421. How do you implement Terraform state locking?

- 422. How do you manage Terraform state securely?
 - 423. How do you protect Terraform state from unauthorized access?
 - 424. How do you use provider aliases for cross-account Terraform deployments?
 - 425. How do you structure Terraform for dev, QA, staging and production?
 - 426. How do you prevent two Terraform pipelines from modifying the same infrastructure simultaneously?
 - 427. How do you detect stale resources created by Terraform automation?
 - 428. How do you safely refactor Terraform modules without recreating production resources?
 - 429. How do you handle Terraform provider version changes?
 - 430. How do you pin Terraform provider versions?
-

25. AWS Fundamentals

- 431. Explain AWS VPC and its components.
- 432. What is a VPC?
- 433. What is a subnet?
- 434. What is the difference between public and private subnets?
- 435. What is an Internet Gateway?
- 436. How do you implement an Internet Gateway?
- 437. How do you configure an Internet Gateway with a route table?
- 438. What is a NAT Gateway?
- 439. What is the purpose of a NAT Gateway?
- 440. How does a NAT Gateway work?
- 441. What is an IAM role and why is it used?
- 442. How does an IAM role differ from an IAM user?
- 443. How do you secure secrets in AWS?
- 444. What AWS services can be used to store secrets?
- 445. What is S3 versioning?

- 446. What is an S3 lifecycle policy?
 - 447. Explain S3 versioning and lifecycle policies.
 - 448. What is the difference between S3 and EBS?
 - 449. What happens when an EC2 instance is stopped versus terminated?
 - 450. What is the difference between EC2, ECS and EKS?
 - 451. When would you use EC2, ECS or EKS?
-

26. AWS EC2

- 452. What is the command to connect to an EC2 instance?
 - 453. What is the exact command to log into an EC2 instance?
 - 454. What type of authentication is used to connect to an EC2 instance?
 - 455. Do you need a password to connect to an EC2 instance?
 - 456. How do you connect to an EC2 instance using SSH?
 - 457. How do you attach an EBS volume to an EC2 instance?
 - 458. How do you attach NFS/EFS to an EC2 instance?
 - 459. What is EBS?
 - 460. What is the difference between EBS and instance store?
 - 461. What happens to EBS volumes when an EC2 instance is stopped?
 - 462. What happens to EBS volumes when an EC2 instance is terminated?
 - 463. How do you troubleshoot an EC2 instance that cannot access the Internet?
 - 464. How do you troubleshoot an EC2 instance that cannot be accessed over SSH?
-

27. AWS VPC / Networking

- 465. What are the dependencies required for an IP-related AWS resource?
- 466. What are the dependencies for an Elastic IP?
- 467. How do you configure a public subnet?
- 468. How do you configure a private subnet?

- 469. How does routing work inside a VPC?
 - 470. How does a private subnet access the Internet?
 - 471. How does a public subnet access the Internet?
 - 472. What is the difference between a Security Group and a Network ACL?
 - 473. What are the differences between Security Groups and NACLs?
 - 474. What is VPC peering?
 - 475. What are the limitations of VPC peering?
 - 476. How do you connect two VPCs?
 - 477. How do you troubleshoot communication between two EC2 instances in different subnets?
 - 478. How do you troubleshoot communication between two VPCs?
 - 479. What is a route table?
 - 480. How does a route table work?
 - 481. How do you troubleshoot a routing issue in AWS?
 - 482. What is the difference between an Internet Gateway and NAT Gateway?
 - 483. What is the difference between NAT Gateway and NAT Instance?
 - 484. An EC2 instance is in a private subnet and cannot use a NAT Gateway. How else can it reach the Internet for updates or downloads?
 - 485. Your EC2 instance in a private subnet needs to download packages without a NAT Gateway. What alternatives exist?
 - 486. How would you securely provide Internet-independent access to AWS services from private subnets?
-

28. AWS Load Balancers / Auto Scaling

- 487. What is the difference between ALB, NLB and CLB?
- 488. How does an ALB work?
- 489. How does an NLB work?
- 490. When would you use an ALB?
- 491. When would you use an NLB?

- 492. How does Auto Scaling work?
 - 493. How does Auto Scaling work with an ALB?
 - 494. How do you configure an Auto Scaling Group?
 - 495. How do you confirm whether Auto Scaling is working as expected?
 - 496. How do you troubleshoot an Auto Scaling Group that is not scaling?
 - 497. How do you configure health checks for an Auto Scaling Group?
 - 498. How do you implement high availability using multiple Availability Zones?
-

29. AWS Cross-Account Access / S3

- 499. You have an application in Account A that needs to access an S3 bucket in Account B. How would you set this up securely?
 - 500. You've got an application running in one AWS account that needs to access an S3 bucket in another account. How would you set that up securely?
 - 501. How do you configure cross-account S3 access?
 - 502. How do IAM roles enable cross-account access?
 - 503. How do you configure an S3 bucket policy for cross-account access?
 - 504. How do you securely access an S3 bucket from an EC2 instance?
 - 505. How do you secure an S3 bucket?
 - 506. What is S3 bucket versioning?
 - 507. What are S3 lifecycle policies?
 - 508. How do you handle accidental deletion of S3 objects?
 - 509. How do you implement S3 backup and retention?
-

30. AWS IAM / Security

- 510. What is an IAM role?
- 511. Why do we use IAM roles?
- 512. What is the difference between an IAM role and IAM user?
- 513. How do you implement least-privilege access?
- 514. How do you secure AWS credentials?

- 515. How do you manage secrets in AWS?
 - 516. What is AWS Secrets Manager?
 - 517. What is AWS Systems Manager Parameter Store?
 - 518. What is the difference between Secrets Manager and Parameter Store?
 - 519. How do you prevent developers from accessing production secrets?
 - 520. How do you enforce least-privilege access for DevOps teams across cloud platforms?
-

31. AWS High Availability / Disaster Recovery

- 521. How do you design multi-zone failover for stateful services with near-zero RPO?
 - 522. What is your approach to testing disaster recovery plans without impacting production traffic?
 - 523. How do you handle DNS-level failures in cloud-based applications?
 - 524. How would you detect and respond to partial regional outages?
 - 525. How do you design a highly available AWS architecture?
 - 526. How do you design multi-AZ applications?
 - 527. How do you design disaster recovery for an AWS application?
 - 528. What is RPO?
 - 529. What is RTO?
 - 530. How do you achieve near-zero RPO?
 - 531. How do you test a disaster recovery strategy?
-

32. Cloud Infrastructure / Automation

- 532. How do you automate provisioning of secure cloud environments for new projects?
- 533. What is your strategy for managing network dependencies such as VPCs, subnets and gateways across environments?
- 534. How do you implement cloud resource tagging policies and enforce them?

- 535. How do you detect and remediate orphaned resources in cloud accounts?
 - 536. What is your strategy for detecting stale or abandoned cloud resources?
 - 537. How do you automate cloud resource cleanup?
 - 538. How do you implement governance across multiple AWS accounts?
 - 539. How do you enforce policies across AWS accounts?
 - 540. How do you implement policy-as-code in cloud environments?
-

33. Monitoring / Prometheus / Grafana

- 541. What is Prometheus?
- 542. What is the configuration file of Prometheus?
- 543. Where can you configure alerting in Prometheus or Grafana?
- 544. I want to configure an alert when disk usage on one server reaches 80%. Where would you configure this?
- 545. How does Prometheus scrape metrics?
- 546. What is a ServiceMonitor?
- 547. How does Prometheus discover Kubernetes targets?
- 548. How do you configure Prometheus to monitor Kubernetes workloads?
- 549. How do you configure Grafana with Prometheus?
- 550. How do you create dashboards in Grafana?
- 551. How do you troubleshoot a missing metric in Grafana?
- 552. How do you troubleshoot Prometheus not scraping a target?
- 553. What is the difference between Prometheus alerting and Grafana alerting?
- 554. How do you configure notification channels for alerts?
- 555. How do you monitor Kubernetes cluster health?
- 556. What metrics do you monitor for Kubernetes nodes?
- 557. What metrics do you monitor for Kubernetes Pods?
- 558. What metrics do you monitor for applications?

34. Logging / Observability / Distributed Tracing

- 559. How would you design a unified logging architecture for services deployed across AWS and Azure?
- 560. What is your approach to identifying cold starts, latency spikes and dropped events in event-driven systems?
- 561. How do you distinguish application bugs from infrastructure issues during incidents?
- 562. What metrics do you prioritize when troubleshooting slow API responses?
- 563. How do you design alerting rules that adapt to business-hour versus off-hour traffic patterns?
- 564. How do you correlate logs, metrics and traces for faster root-cause analysis?
- 565. How do you debug a production issue when logs show no errors?
- 566. How do you identify root cause when metrics look normal but users complain?
- 567. What signals do you trust most during a live production outage?
- 568. How do you detect silent failures?
- 569. How do you troubleshoot latency spikes?
- 570. How do you troubleshoot dropped events in event-driven systems?
- 571. How do you implement distributed tracing?
- 572. How do you correlate a request across multiple microservices?
- 573. How do you use trace IDs in application logs?
- 574. How do you distinguish application latency from infrastructure latency?

35. Security / DevSecOps

- 575. How do you manage secrets and credentials across CI/CD pipelines and runtime environments?
- 576. What is your approach to integrating vulnerability scanning into release pipelines?

- 577. How do you enforce least-privilege access for DevOps teams across cloud platforms?
 - 578. How do you ensure audit readiness for enterprise compliance requirements?
 - 579. How do you integrate Trivy into a CI/CD pipeline?
 - 580. How do you integrate Snyk into a CI/CD pipeline?
 - 581. How do you integrate SonarQube into Jenkins?
 - 582. What is Software Composition Analysis?
 - 583. How do you scan application dependencies for vulnerabilities?
 - 584. How do you scan Docker images for vulnerabilities?
 - 585. What is image signing?
 - 586. How does Cosign work?
 - 587. How do you prevent unsigned images from being deployed?
 - 588. What is OPA?
 - 589. What is Conftest?
 - 590. What is Kyverno?
 - 591. How do you enforce Kubernetes admission policies?
 - 592. How do you implement policy-as-code?
 - 593. How do you prevent secrets from being committed to Git?
 - 594. How do you use Gitleaks?
 - 595. How do you implement security gates in CI/CD?
-

36. Cloud Cost / Governance

- 596. Your company's cloud costs are increasing rapidly. How would you approach cost optimization without impacting performance?
- 597. If your team is seeing a spike in cloud costs, how would you determine why and reduce the cost without hurting performance?
- 598. How do you identify which AWS services are causing a cost increase?
- 599. How do you optimize EC2 costs?

- 600. How do you optimize EKS costs?
 - 601. How do you optimize S3 costs?
 - 602. How do you identify unused cloud resources?
 - 603. How do you identify orphaned resources?
 - 604. How do you implement cloud resource tagging?
 - 605. How do you enforce mandatory tags?
 - 606. How do you implement cost allocation across teams?
 - 607. How do you monitor cloud cost trends?
-

37. AWS Geolocation / DNS

- 608. How would you set up geolocation-based routing using AWS services?
 - 609. Say you want to serve users in different countries using AWS. How would you route traffic based on user location?
 - 610. What is Route 53?
 - 611. Is it possible to purchase a domain using Route 53?
 - 612. What are Route 53 routing policies?
 - 613. What is geolocation routing?
 - 614. What is latency-based routing?
 - 615. What is failover routing?
 - 616. How do you implement DNS failover?
 - 617. How do you troubleshoot DNS-level failures?
-

38. Advanced CI/CD Reliability

- 618. How do you design a controlled deployment flow where every release must pass environment-specific validation before progressing?
- 619. What is your approach to artifact promotion across dev → QA → staging → production without rebuilding?
- 620. How would you ensure pipeline self-healing when a shared agent becomes unresponsive mid-deployment?

- 621. How do you handle parallel infrastructure changes in a monorepo environment to avoid conflicting pipeline triggers?
 - 622. How do you design a CI/CD pipeline that supports multiple release calendars for different business units?
 - 623. What is your strategy for enforcing branch protection and release gates in large repositories?
 - 624. How do you handle failed deployments when rollback scripts are not available?
 - 625. How do you implement artifact versioning and promotion across environments?
 - 626. How do you design a CI/CD pipeline where failed deployments automatically freeze further releases until the root cause is identified?
 - 627. How would you implement deployment verification tests for systems that cannot tolerate downtime?
 - 628. How would you manage version pinning across multiple dependent microservices?
 - 629. How do you ensure CI pipelines remain deterministic when builds depend on external services?
-

39. Kubernetes Runtime Stability / Platform Reliability

- 630. How do you debug sporadic Pod restarts when logs show no application errors?
- 631. What steps do you take if kubelet health degrades on only a subset of nodes?
- 632. How do you protect critical workloads from noisy-neighbor issues in shared clusters?
- 633. How would you safely roll out kernel or node-image upgrades in production clusters?
- 634. A StatefulSet Pod is failing only on a specific node. How do you isolate whether it is a node issue, storage issue or scheduling conflict?
- 635. How do you safely rotate TLS certificates used by Ingress controllers without service disruption?

- 636. What is your strategy to enforce namespace-level resource isolation in a multi-team Kubernetes cluster?
 - 637. How would you validate the impact of a new NetworkPolicy before applying it cluster-wide?
 - 638. How do you troubleshoot uneven Pod traffic distribution behind a Service?
 - 639. How do you troubleshoot Pod evictions?
 - 640. How do you troubleshoot ImagePullBackOff?
 - 641. How do you troubleshoot 503 errors?
 - 642. How do you troubleshoot multiple simultaneous Kubernetes failures?
-

40. Cloud Architecture / High Availability

- 643. How do you design multi-zone failover for stateful services with near-zero RPO?
 - 644. What is your approach to testing disaster recovery plans without impacting production traffic?
 - 645. How do you handle DNS-level failures in cloud-based applications?
 - 646. How would you detect and respond to partial regional outages?
 - 647. How do you design an application for high availability across Availability Zones?
 - 648. How do you design stateful services for disaster recovery?
 - 649. How do you handle regional failures?
 - 650. How do you validate a disaster recovery plan?
 - 651. How do you perform a DR drill without impacting production?
 - 652. How do you design backup and restore strategies for critical applications?
-

41. Production Incident / Troubleshooting Scenarios

- 653. How do you debug a production issue when logs show no errors?
- 654. What steps do you follow when application latency suddenly increases?

655. How do you identify the root cause when metrics look normal but users complain?
656. What do you check first when a Kubernetes Pod is running but traffic fails?
657. How do you troubleshoot intermittent failures that cannot be reproduced?
658. How do you verify whether a deployment actually changed runtime behavior?
659. What do you analyze when database connections are exhausted but CPU is low?
660. How do you approach an incident involving multiple dependent services?
661. What signals do you trust most during a live production outage?
662. How do you confirm whether autoscaling is working as expected?
663. How do you debug a pipeline that succeeds but produces a faulty release?
664. What would you investigate if a rollback does not resolve the issue?
665. How do you detect configuration drift in production systems?
666. What checks do you perform before declaring an incident resolved?
667. How do you prioritize actions during a high-severity outage?
668. How do you distinguish between application bugs and infrastructure issues during incidents?
669. What metrics do you prioritize when troubleshooting slow API responses?
670. How do you correlate logs, metrics and traces for faster root-cause analysis?
671. How do you debug a production issue when the application logs appear normal?
672. How do you troubleshoot a problem that cannot be reproduced?
673. How do you determine whether a deployment caused a production issue?
674. How do you determine whether an issue is caused by infrastructure or application code?
675. How do you troubleshoot database connection exhaustion when CPU and memory look normal?

42. Large Kubernetes Production Scenario

676. A critical production Kubernetes cluster is experiencing multiple issues. Pods are stuck in ImagePullBackOff, some Pods are being evicted, and users are reporting 503 errors. What troubleshooting process will you follow, and how can you avoid this in the future?
677. You're on-call. A production Kubernetes cluster is a mess: Pods aren't pulling images, some are getting evicted, and users are seeing errors. How do you troubleshoot this, and how would you prevent it next time?

43. Behavioral Round

678. How do you handle a situation where you're asked to work on a technology you have no experience with?
679. What would you do if you were asked to work on a tool or technology you've never used before?
680. Describe a time when you had to work with tight deadlines and limited resources.
681. Can you share a time when you had to deliver something quickly but didn't have all the time or resources you needed? How did you manage it?
682. Tell me about a mistake you made in production and how you handled it.
683. Tell me about a mistake or outage you were involved in. How did you respond, and what did you learn from it?
684. Describe the most challenging technical problem you've solved in your career.
685. What's the hardest technical issue you've faced so far? How did you go about solving it?
686. How would you convince stakeholders to adopt a new technology or process?
687. If you wanted to convince your team or manager to adopt a new tool or process, how would you make your case?
688. Tell me about a time when you had to learn a new tool quickly to solve a business problem.

- 689. Describe a time when you had to troubleshoot a production issue under pressure.
 - 690. Tell me about a time when you disagreed with a technical decision.
 - 691. Tell me about a time when you automated a manual process.
 - 692. Tell me about a time when you improved a deployment process.
 - 693. Tell me about a time when you handled a production incident.
 - 694. Tell me about a time when you made a mistake and what you learned from it.
-

44. Additional Scenario-Based Questions

- 695. How do you design a deployment flow where every environment has its own validation?
- 696. How do you promote a single immutable artifact across multiple environments?
- 697. How do you prevent rebuilding an application when promoting it from QA to production?
- 698. How do you make a deployment automatically stop when validation fails?
- 699. How do you implement deployment freeze after a production failure?
- 700. How do you implement automated deployment verification?
- 701. How do you safely deploy applications that cannot tolerate downtime?
- 702. How do you manage multiple microservice versions and dependencies?
- 703. How do you prevent external dependencies from making CI builds non-deterministic?
- 704. How do you handle an unavailable Jenkins agent during a deployment?
- 705. How do you handle concurrent infrastructure deployments?
- 706. How do you prevent multiple pipelines from modifying the same Terraform state?
- 707. How do you handle infrastructure changes in a monorepo?
- 708. How do you implement release gates?
- 709. How do you implement artifact immutability?

- 710. How do you implement rollback when no rollback script exists?
 - 711. How do you determine whether rollback actually restored the previous application behavior?
-

45. More AWS / Infrastructure Interview Questions

- 712. How do you provision a secure AWS environment using Terraform?
 - 713. How do you provision networking using Terraform?
 - 714. How do you create an AWS VPC using Terraform?
 - 715. How do you create public and private subnets using Terraform?
 - 716. How do you create an Internet Gateway using Terraform?
 - 717. How do you create NAT Gateway infrastructure using Terraform?
 - 718. How do you create an ALB using Terraform?
 - 719. How do you create an Auto Scaling Group using Terraform?
 - 720. How do you create an EKS cluster using Terraform?
 - 721. How do you manage shared infrastructure across multiple Terraform environments?
 - 722. How do you manage multiple AWS accounts using Terraform?
 - 723. How do you manage multiple AWS regions using Terraform?
 - 724. How do you reuse Terraform modules across multiple accounts?
 - 725. How do you detect resources created manually outside Terraform?
 - 726. How do you bring manually created resources under Terraform management?
 - 727. How do you safely modify Terraform-managed production infrastructure?
 - 728. How do you review Terraform changes before applying them?
-

46. Additional Kubernetes Interview Questions

- 729. How does the Kubernetes scheduler decide where a Pod runs?
- 730. What are node selectors?
- 731. What are node affinity and anti-affinity?

- 732. What are taints and tolerations?
 - 733. How do taints and tolerations affect Pod scheduling?
 - 734. How do you ensure Pods are distributed across multiple nodes?
 - 735. How do you ensure Pods are distributed across multiple Availability Zones?
 - 736. What are PodDisruptionBudgets?
 - 737. Why are PodDisruptionBudgets used?
 - 738. How do you perform a controlled node drain?
 - 739. How do you prevent critical Pods from being evicted?
 - 740. How do readiness probes affect traffic?
 - 741. How do liveness probes affect Pod restarts?
 - 742. What happens if a readiness probe fails?
 - 743. What happens if a liveness probe fails?
 - 744. How do you troubleshoot failing readiness probes?
 - 745. How do you troubleshoot failing liveness probes?
 - 746. How do you troubleshoot Kubernetes DNS?
 - 747. How do you troubleshoot CoreDNS?
 - 748. How do you troubleshoot Service-to-Pod communication?
 - 749. How do you troubleshoot Pod-to-Pod communication?
 - 750. How do you troubleshoot external-to-Pod communication?
-

47. Questions Appearing as Repeated / Near-Repeated Versions

I am **intentionally retaining these**, because you specifically said not to remove questions just because they are similar.

- 751. What is the difference between Deployment and StatefulSet in Kubernetes?
- 752. What is the difference between Deployment, StatefulSet and DaemonSet?
- 753. How does Kubernetes perform rolling updates?

- 754. How do you configure a rolling update in Kubernetes?
- 755. What is the difference between ConfigMap and Secret?
- 756. What is a ConfigMap, and how is it different from a Secret?
- 757. How do Kubernetes Services find Pods?
- 758. How are Kubernetes Services linked to Pods?
- 759. How does one namespace Pod find another namespace?
- 760. How do Pods communicate across namespaces?
- 761. How do you troubleshoot a Pod stuck in CrashLoopBackOff?
- 762. How do you debug sporadic Pod restarts when logs show no application errors?
- 763. How do you debug a production issue when logs show no errors?
- 764. How do you handle Terraform state corruption?
- 765. Suppose your Terraform state file gets corrupted or out of sync. What steps would you take?
- 766. How do you import an existing AWS VPC into Terraform?
- 767. You have an existing AWS VPC created manually. How do you bring it under Terraform?
- 768. How do you handle an S3 bucket policy manually changed outside Terraform?
- 769. Someone manually added a policy to an S3 bucket created by Terraform. How do you handle the drift?
- 770. What is the difference between Security Groups and NACLs?
- 771. Difference between Security Groups and NACLs.
- 772. How do you implement cross-account AWS access?
- 773. An application in Account A needs to access S3 in Account B. How do you configure it?
- 774. How do you implement blue-green deployment in Kubernetes?
- 775. How would you roll out blue-green deployment in Kubernetes?
- 776. How do you securely manage Terraform secrets?
- 777. How do you manage secrets in Terraform without hardcoding them?

- 778. What is the difference between COPY and ADD?
 - 779. What's the practical difference between COPY and ADD in Docker?
 - 780. How do you pass secrets to a PHP Docker container?
 - 781. How would you handle secrets in a Docker container for a PHP application connecting to MySQL?
 - 782. How do you implement NetworkPolicies in Kubernetes?
 - 783. How do you enforce rules controlling Pod-to-Pod communication?
 - 784. How do you optimize cloud costs?
 - 785. Your cloud costs are increasing rapidly. How would you optimize them without impacting performance?
 - 786. How do you implement geolocation routing in AWS?
 - 787. How would you route users based on their location using AWS?
 - 788. How do you troubleshoot ImagePullBackOff?
 - 789. A production Kubernetes cluster has ImagePullBackOff, Pod evictions and 503 errors. How do you troubleshoot it?
-

48. Specific Practical Questions From the Screenshots

- 790. I have an index.html on GitHub. I need to push it into AWS. The requirement is one load balancer, and we provide index.html in the GitHub repository along with username and token. How do you achieve this?
- 791. What are the dependencies/resources for the IP?
- 792. What is the command to connect to an EC2 instance?
- 793. How do you implement an Internet Gateway?
- 794. How do you configure an Internet Gateway to a route table?
- 795. What is the exact command to log into an instance and what kind of authentication do you use?
- 796. Do you need to provide a password to connect to an EC2 instance?
- 797. In which path will you create index.html on the server?
- 798. Explain the most common Linux commands.
- 799. How do you check one running process?

- 800. How do you list running processes?
 - 801. How do you find a particular process ID and kill it?
 - 802. Is it possible to find and kill a process using a single command?
 - 803. How do you check disk usage?
 - 804. How do you find free memory?
 - 805. How do you archive and compress a directory in Linux?
 - 806. What does `chmod 755` mean?
 - 807. If I provide `chmod 755`, what exactly will happen?
 - 808. What is shown by `ls`?
 - 809. How do you list all SSH users in Linux?
 - 810. If an Apache server is running and you need to check its logs, where do you check?
 - 811. In which directory are Apache logs stored?
 - 812. What kind of logs can you see in `/var/log`?
 - 813. What is the purpose of the `grep` command?
 - 814. If I need to `grep` `linux` from `linux.txt`, and I expect the keyword `linux`, which flag will you use?
-

49. Jenkins / Application Deployment Practical Questions

- 815. If you launched the service and installed Jenkins but Jenkins is not accessible, what might be the reason?
- 816. What are the commands called `git cherry-pick`, `git config` and `git rebase`?
- 817. How do you install Jenkins plugins?
- 818. What are the steps to install plugins?
- 819. What are pipeline stages?
- 820. Can you write Jenkins pipeline stages?
- 821. What is the Maven lifecycle?
- 822. What is the `mvn clean install` command?
- 823. What is a WAR file?

- 824. Where do you store the artifact?
 - 825. What is the location from which you get the artifact file?
 - 826. How do you integrate SonarQube with Jenkins?
 - 827. What is CI/CD? Explain it.
 - 828. What are CMD and ENTRYPOINT?
 - 829. What are COPY and ADD?
 - 830. Can you write a Kubernetes Deployment manifest?
 - 831. What happens after running Node.js?
 - 832. How do you run a specific task using Ansible?
 - 833. What S3 buckets have you used?
-

50. Advanced Platform / Reliability Questions

- 834. How do you design a controlled deployment flow with environment-specific validation?
- 835. How do you implement artifact promotion without rebuilding?
- 836. How do you make pipelines self-healing?
- 837. How do you avoid conflicting pipeline triggers in a monorepo?
- 838. How do you isolate StatefulSet failures to node, storage or scheduling?
- 839. How do you rotate TLS certificates without downtime?
- 840. How do you enforce namespace-level isolation?
- 841. How do you validate NetworkPolicies before cluster-wide rollout?
- 842. How do you structure multi-account, multi-region Terraform?
- 843. How do you detect stale resources?
- 844. How do you migrate from scripts to Terraform?
- 845. How do you create unified logging across AWS and Azure?
- 846. How do you identify cold starts and dropped events?
- 847. How do you distinguish application bugs from infrastructure problems?
- 848. How do you troubleshoot slow API responses?

- 849. How do you design adaptive alerting?
 - 850. How do you correlate logs, metrics and traces?
 - 851. How do you manage secrets across CI/CD and runtime?
 - 852. How do you integrate vulnerability scanning into release pipelines?
 - 853. How do you enforce least privilege?
 - 854. How do you maintain audit readiness?
-

51. Important Topics Appearing in the Screenshots but Not Written as Questions

These were shown as **topics/headings rather than complete questions**, so I am keeping them separately instead of pretending they were questions:

- S3 versioning and lifecycle policy
- Dockerfile vs Docker Compose
- Docker volumes
- Terraform
- AWS VPC
- AWS IAM
- Jenkins / CI-CD
- Kubernetes Components
- Kubernetes & Platform Reliability
- Infrastructure as Code
- Logging, Monitoring & Distributed Tracing
- Observability & Production Troubleshooting
- Security & Compliance Automation
- Cloud Infrastructure & Automation
- Cloud Architecture & High Availability
- Kubernetes & Application Operations
- GitLab CI/CD
- GitHub Actions

- Docker
- Bash Scripting
- Python Scripting
- AWS
- Git
- Kubernetes
- ArgoCD / GitOps
- DevSecOps
- Behavioral Round